

the science of deep specification

Verification of interacting systems in VST

a.k.a Towards a verified web server, part II

Lennart Beringer,
William Mansky



Joachim Breitner, **Nicolas Koh**,
Yao Li, **Yishuai Li**, Benjamin C.
Pierce, **Li-Yao Xia**, Steve
Zdancewic



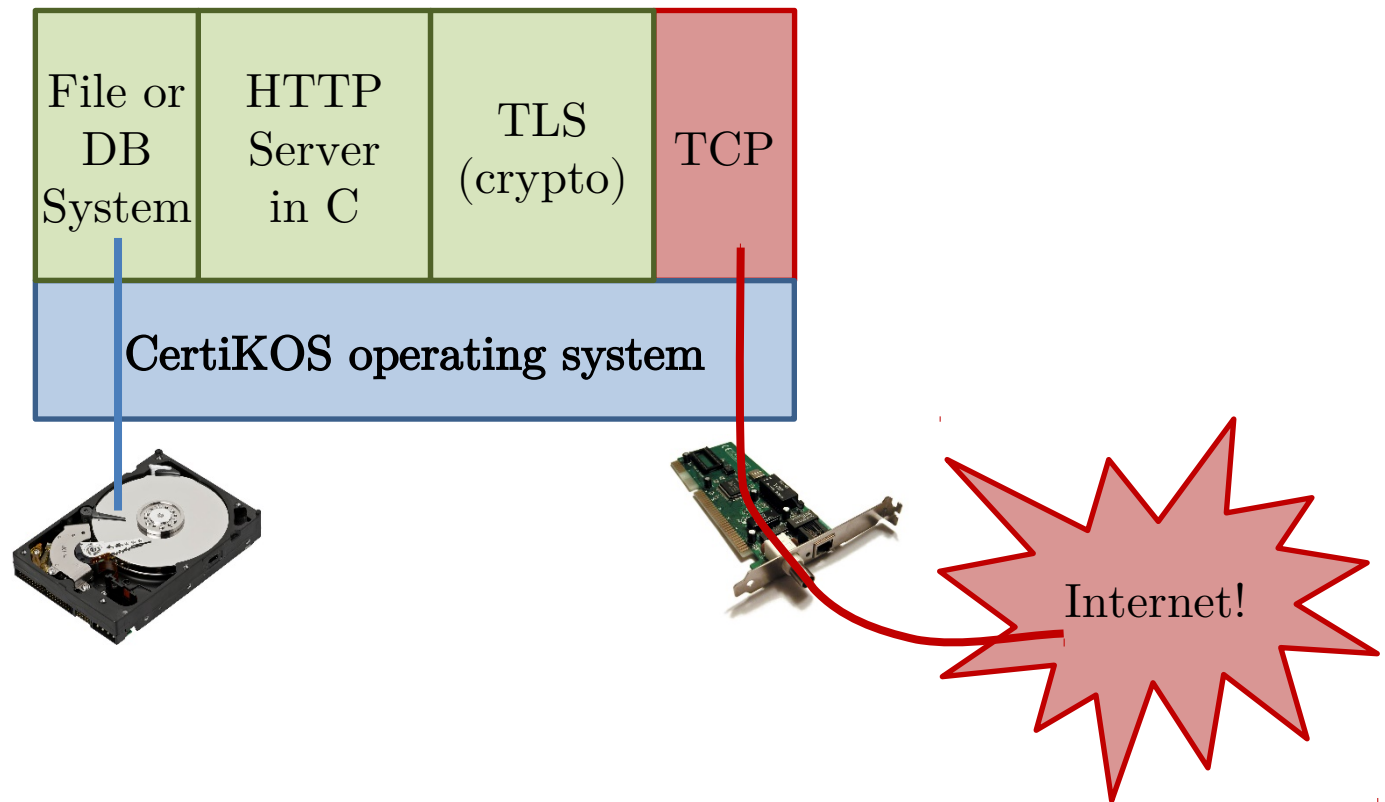
Wolf Honore




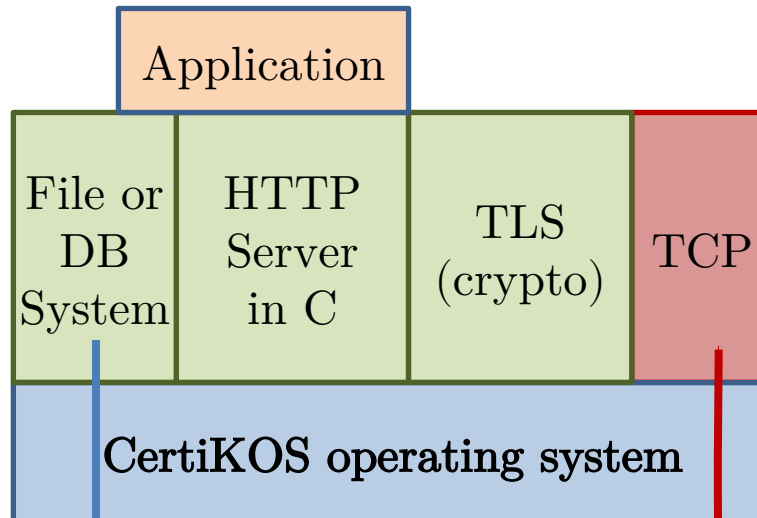

Demo vehicle for DeepSpec Techniques




DeepSpec Web Server



Server as a library




Challenges



- complexity / code size
- interfaces between components at best informally / intuitively specified
- specifications based on internal view of implementations unconnected to specification of externally visible behavior (messages sent on network)
- ...

Challenges

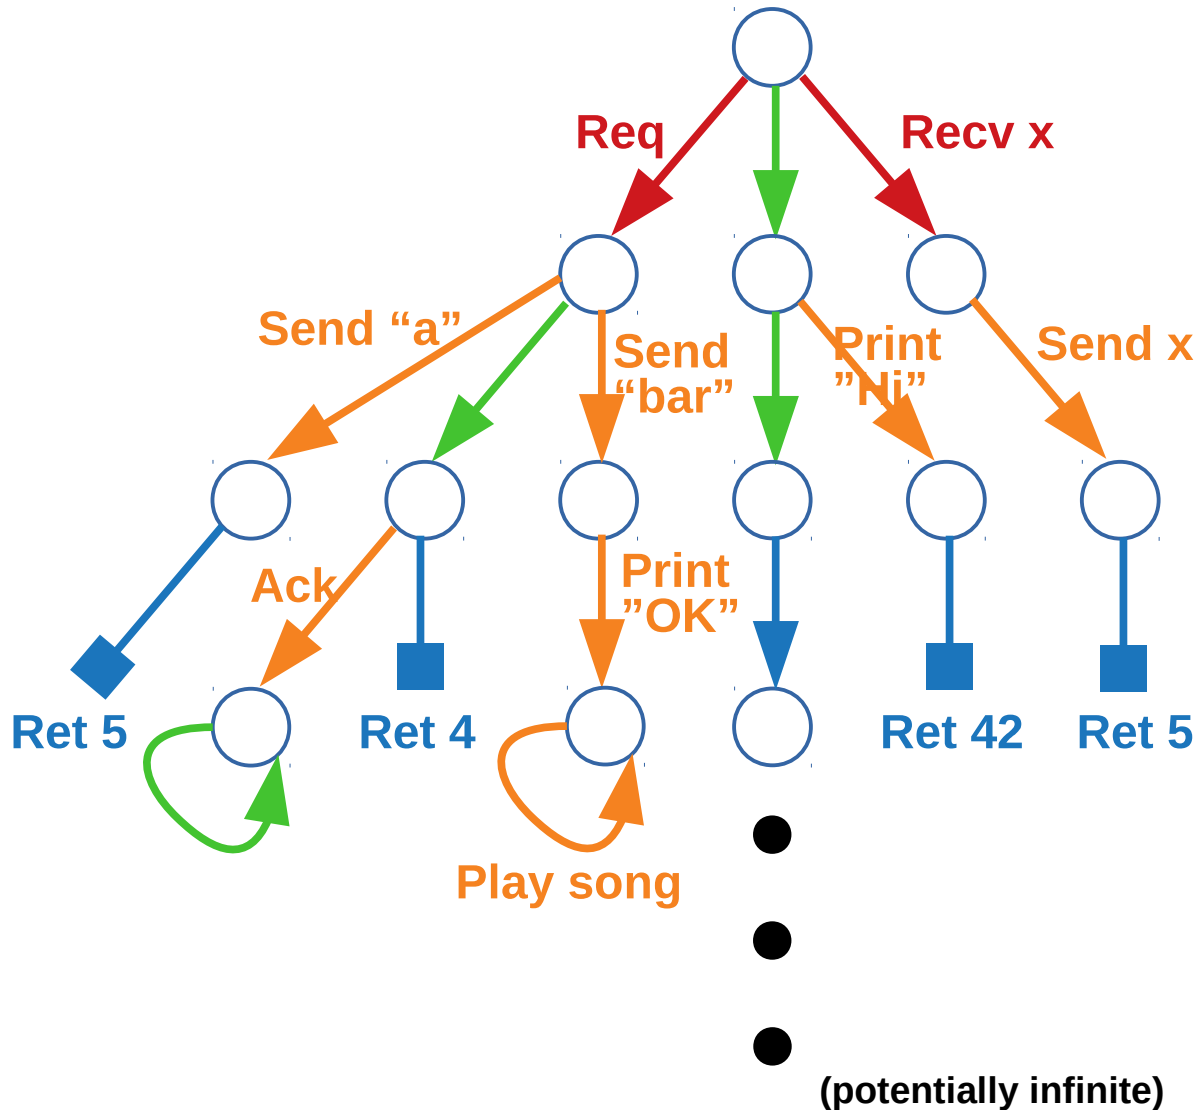


- complexity / code size
- interfaces between components at best informally / intuitively specified
- specifications based on internal view of implementations unconnected to specification of externally visible behavior (messages sent on network)
- ...

Today: specification + verification of effectful code in VST

- OS primitives for sending and receiving network packets
-
- link model-level specification of interactions in Gallina / Quickchick to code-level specs in VST

Interaction trees



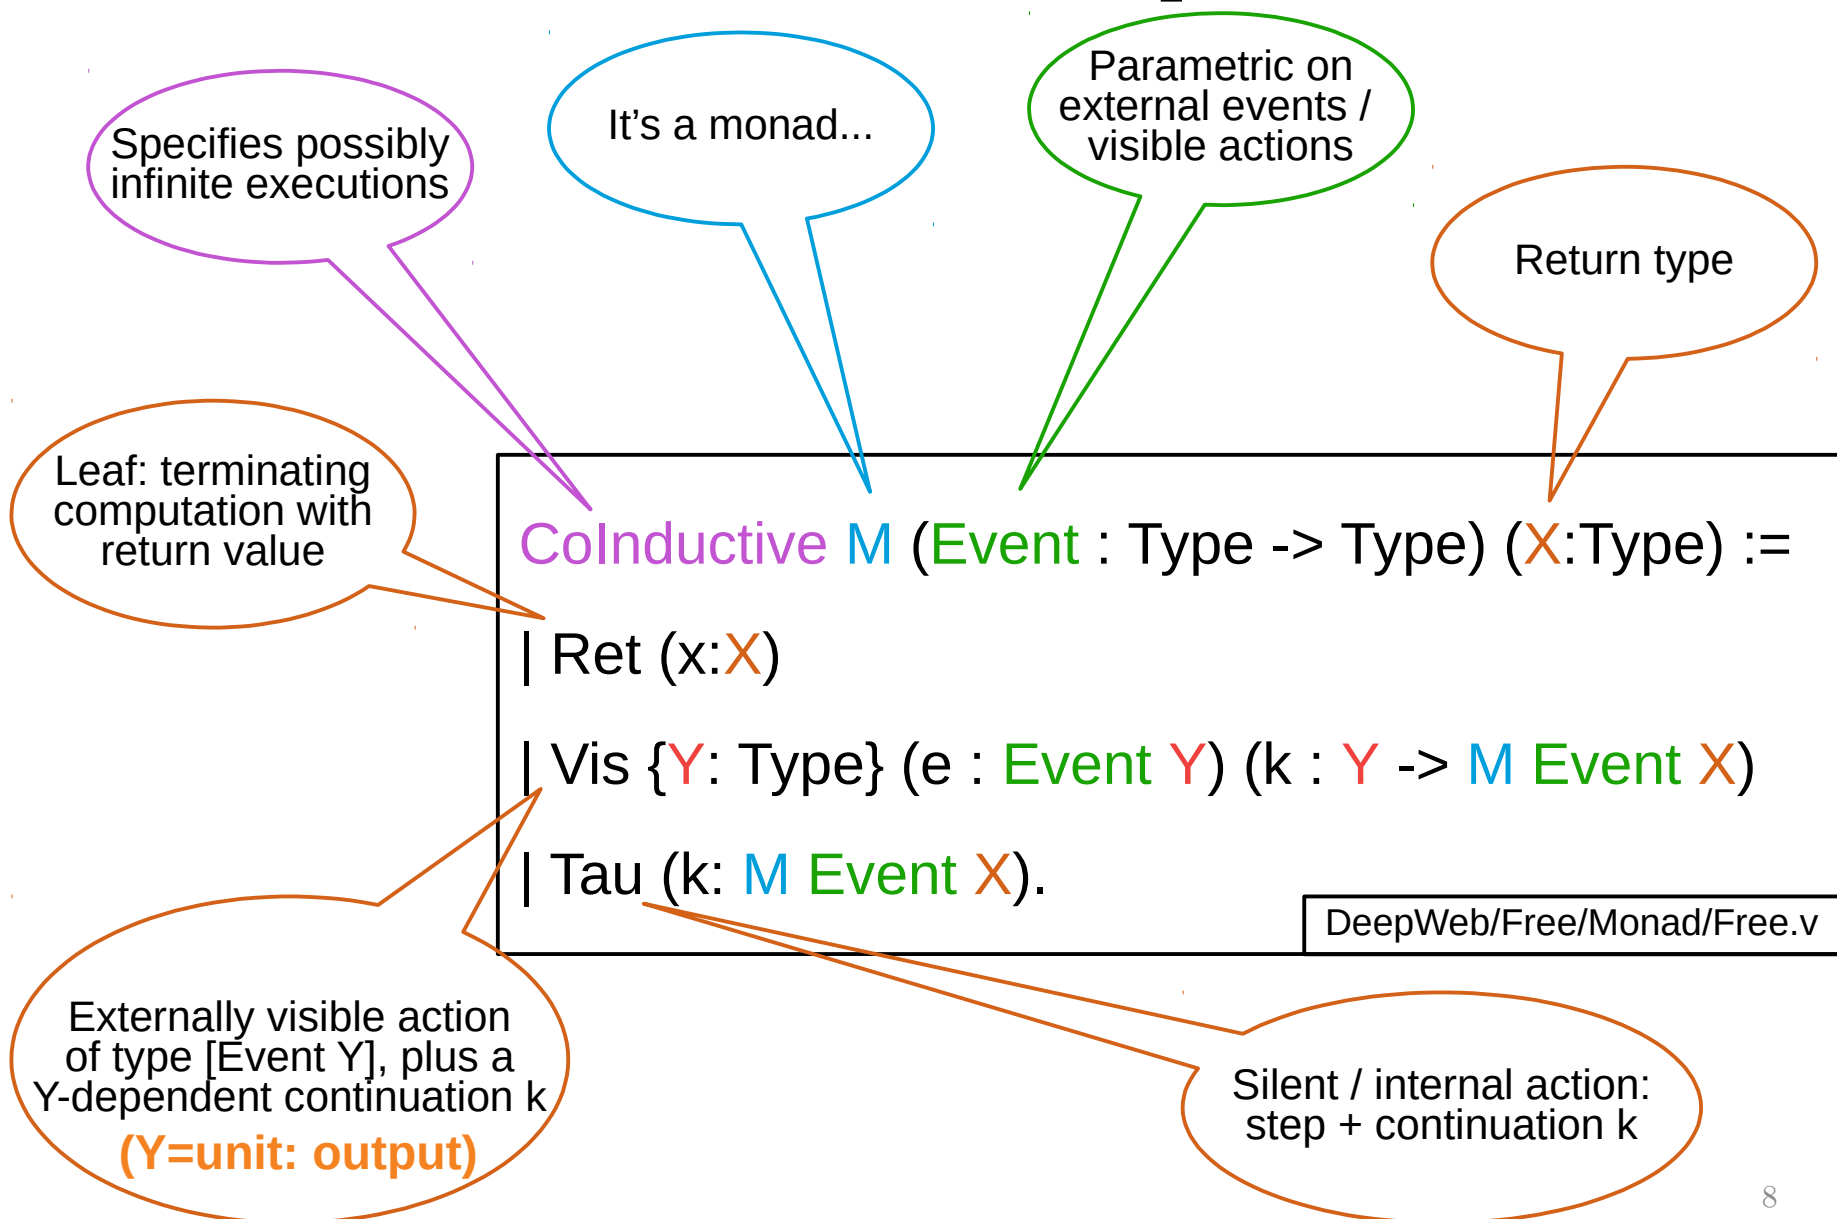
Input actions

Output actions

Silent (tau) actions

Terminal action

Interaction trees in Coq: Definition



Interaction: Properties + Constructions

```
CoInductive M (Event : Type -> Type) (X:Type) :=  
  | Ret (x:X)  
  | Vis {Y: Type} (e : Event Y) (k : Y -> M Event X)  
  | Tau (k: M Event X).
```

Monad structure: Ret +

```
Definition bindM {E X Y} (s: M E X) (t: X -> M E Y) : M E Y := ...  
(notation y ← s ;; t y).
```

Single-event tree: Definition liftE E X (e : E X) : M E X := Vis e (fun x => Ret x).

Discard result: Definition ignore {E X}: M E X -> M E unit := ...

Silent infinite loop: CoFixpoint spin {E X} : M E X := Tau spin.

Interacting inf. loop: CoFixpoint forever {E X} (x : M E X) : M E void := x;; forever x.

Interaction: Properties + Constructions

Questions:

1. How to link interaction trees to C code in VST ?
2. How to specify OS primitives for sending and receiving messages ?

```
Coinductive M (Event : Type -> Type) (X:Type) :=  
  | Ret (x:X)  
  | Vis {Y: Type} (e : Event Y) (k : Y -> M Event Y)  
  | Tau (k: M Event X).
```

```
Definition bindM {E X Y} (s: M E X) (t: X -> M E Y) : M E Y := ...  
(notation y ← s ;; t y).
```

Single-event tree: `Definition liftE E X (e : E X) : M E X := Vis e (fun x => Ret x).`

Discard result: `Definition ignore {E X}: M E X -> M E unit := ...`

Silent infinite loop: `CoFixpoint spin {E X} : M E X := Tau spin.`

Interacting inf. loop: `CoFixpoint forever {E X} (x : M E X) : M E void := x;; forever x.`

Interaction: Properties + Constructions

Questions:

1. How to link interaction trees to C code in VST ?
2. How to specify OS primitives for sending and receiving messages ?

Today:

1. Motivate ITREE separation logic predicates
2. Two examples: character-IO, swap server

Interaction trees in VST

(maintained in SEP clause)



- embed tree **t** in a novel **spatial** predicate **ITREE (t)**.

```
PRE ...  
  PROP(...) LOCAL (...) SEP(...; ITREE (r ← x;; k r))  
...  
POST ...  
  EX v,  
  PROP(v =...) LOCAL (...) SEP(...; ITREE (k v))
```

- Forward direction (“consuming” style)
 - cf. actions in process calculi
 - enables referring to input in continuation via binding
 - compatibility with partial-correctness (eg infinite loops)

Interaction trees in VST: semantic foundation

VST's supports ghost state, allowing users to specify
ghost variables / locations that can hold arbitrary Coq values

VST/veric/juicy_extspec.v

Definition `has_ext` {Z:Type} (ora:Z):mpred :=

Informal reading:

- ghost location 0 stores value ora:Z
- shared ownership: location can be
 - **read** by code **and** the external world (OS)
 - **modified only** by external world

Interaction trees in VST: semantic foundation

VST's supports ghost state, allowing users to specify
ghost variables / locations that can hold arbitrary Coq values

VST/veric/juicy_extspec.v

Definition `has_ext` {Z:Type} (ora:Z):mpred :=

Informal reading:

- ghost location 0 stores value ora:Z
- shared ownership: location can be
 - **read** by code **and** the external world (OS)
 - **modified only** by external world

First idea: use `has_ext` directly:

Definition `ITREE` (t: M ...):mpred := `has_ext t`.

- program and external world agree on current “protocol stage”, and only OS can change this
- proof rules essentially embed external function call effects as `ltree` reductions the ghost state

Example: getchar/putchar

1. Definition of I-Trees for read / write interactions

(* The type of actions: reading yields an int, writing yields a unit *)

Inductive IO_event : Type -> Type :=

| ERead : IO_event int

| EWrite (c : int) : IO_event unit.

(* Define shorthands for each of the two operations*)

Definition read : M IO_event int := embed ERead.

Definition write (c : int) : M IO_event unit := embed (EWrite c).

(*Instantiate the I-Tree monad M – overall return type is unit*)

Definition IO_itree := M IO_event unit.

I-Trees in VST: getchar/putchar

2. VST axioms for OS / library functions

Definition `putchar_spec` :=
DECLARE `_putchar`

WITH `c` : int, `k` : unit -> IO_itree

PRE [`_c` OF tint]

PROP ()
LOCAL (temp `_c` (Vint `c`))
SEP (**ITREE** (`write c` ;; `k`))

POST [tint]

PROP ()
LOCAL (temp `ret_temp` (Vint `c`))
SEP (**ITREE** `k`).

Definition `getchar_spec` :=
DECLARE `_getchar`

WITH `k` : int -> IO_itree

PRE []

PROP ()
LOCAL ()
SEP (**ITREE** (`r` <- `read` ;; `k r`))

POST [tint]

EX `i` : int,
PROP ($-2^7 \leq \text{Int.signed } i < 2^7$)
LOCAL (temp `ret_temp` (Vint `i`))
SEP (**ITREE** (`k i`)).

Caution: these specs don't model error cases

I-Trees in VST: getchar/putchar

3. Write some C program: charIO/IO/io.c

```
void print_intr(uint l)
{
    unsigned int q,r;
    if (l!=0) {
        q=l/10u;
        r=l%10u;
        print_intr(q);
        putchar(r+'0');
    }
}
```

```
void print_int(uint i)
{
    if (i==0)
        putchar('0');
    else print_intr(i);
}
```

```
int main(void) {
    unsigned int n, d; char c;

    n=0;
    c=getchar();
    while (n<1000) {
        d = ((unsigned)c)-(unsigned)'0';
        if (d>=10) break;
        n+=d;
        print_int(n);
        putchar('\n');
        c=getchar();
    }
    Return 0;
}
```

I-Trees in VST: getchar/putchar

3. Write some C program: charIO/IO/io.c

```
void print_intr(uint l)
{
    unsigned int q,r;
    if (l!=0) {
        q=l/10u;
        r=l%10u;
        print_intr(q);
        putchar(r+'0');
    }
}
```

```
void print_int(uint i)
{
    if (i==0)
        putchar('0');
    else print_intr(i);
}
```

```
int main(void) {
    unsigned int n, d; char c;

    n=0;
    c=getchar();
    while (n<1000) {
        d = ((unsigned)c)-(unsigned)'0';
        if (d>=10) break;
        n+=d;
        print_int(n);
        putchar('\n');
        c=getchar();
    }
    Return 0;
}
```

4. Write some informal spec

- reads a sequence of characters, each in the range '0'..'9';
- after each one (and before the next one) prints a decimal representation of their cumulative sum, followed by a newline;
- when the cumulative sum ≥ 1000 , or when EOF is reached, stops.

I-Trees in VST: getchar/putchar

5. Translate the informal spec into Coq / ITrees

```
// Itree generating write ops for all elements of a list
Fixpoint write_list (l: list int) : IO_itree :=
match l with
| nil => Ret tt
| c :: rest => write c ;; write_list rest
end.
```

```
Program Fixpoint chars_of_nat n { measure n } : list int := ...
```

```
Definition chars_of_Z z := chars_of_nat (Z.to_nat z).
```

```
// the function computed by print_intr
Program Fixpoint intr n { measure n } : list int :=
match n with
| 0 => []
| _ => intr (n / 10) ++ [Int.repr (Z.of_nat (n mod 10) + char0)]
end.
```

I-Trees in VST: getchar/putchar

6. Embed Itrees in VST specs

```
void print_intr(uint i)
{
  unsigned int q,r;
  if (i!=0) {
    q=i/10u;
    r=i%10u;
    print_intr(q);
    putchar(r+'0');
  }
}
```

```
Fixpoint write_list (l: list int) :=
match l with
| nil => Ret tt
| c :: t => write c ;; write_list t
end.
```

Program Fixpoint intr n := ...

Definition print_intr_spec :=

DECLARE _print_intr

WITH i : Z, tr : IO_itree

PRE [_i OF tuint]

PROP (0 <= i <= Int.max_unsigned)

LOCAL (temp _i (Vint (Int.repr i)))

SEP (ITREE (write_list (intr (Z.to_nat i)) ;; tr))

POST [tvoid]

PROP ()

LOCAL ()

SEP (ITREE tr).

Similarly for print_int function

I-Trees in VST: getchar/putchar

6. Embed Itrees in VST specs

```
int main(void) {  
  unsigned int n, d; char c;  
  
  n=0;  
  c=getchar();  
  while (n<1000) {  
    d = ((unsigned)c)-(unsigned)'0';  
    if (d>=10) break;  
    n+=d;  
    print_int(n);  
    putchar('\n');  
    c=getchar();  
  }  
  Return 0;  
}
```

```
CoFixpoint read_sum n d : IO_itree :=  
  if zlt n 1000 then  
    if zlt d 10 then  
      write_list (chars_of_Z (n + d));;  
      write (Int.repr newline);;  
      c <- read;;  
      read_sum (n + d) (Int.unsigned c - char0)  
    else ret tt else ret tt.
```

```
Definition main_itree := c <- read;; read_sum 0  
(Int.unsigned c - char0).
```

I-Trees in VST: getchar/putchar

6. Embed Itrees in VST specs

```
int main(void) {  
  unsigned int n, d; char c;  
  
  n=0;  
  c=getchar();  
  while (n<1000) {  
    d = ((unsigned)c)-(unsigned)'0';  
    if (d>=10) break;  
    n+=d;  
    print_int(n);  
    putchar('\n');  
    c=getchar();  
  }  
  Return 0;  
}
```

CoFixpoint??? Program terminates
after 1000 iterations!

```
CoFixpoint read_sum n d : IO_itree :=  
  if zlt n 1000 then  
    if zlt d 10 then  
      write_list (chars_of_Z (n + d));;  
      write (Int.repr newline);;  
      c <- read;;  
      read_sum (n + d) (Int.unsigned c - char0)  
    else ret tt else ret tt.
```

But really finite !

```
Definition main_itree := c <- read;; read_sum 0  
(Int.unsigned c - char0).
```

I-Trees in VST: getchar/putchar

6. Embed Itrees in VST specs

```
int main(void) {  
  unsigned int n, d; char c;  
  
  n=0;  
  c=getchar();  
  while (n<1000) {  
    d = ((unsigned)c)-(unsigned)'0';  
    if (d>=10) break;  
    n+=d;  
    print_int(n);  
    putchar('\n');  
    c=getchar();  
  }  
  Return 0;  
}
```

CoFixpoint??? Program terminates
after 1000 iterations!

```
CoFixpoint read_sum n d : IO_itree :=  
  if zlt n 1000 then  
    if zlt d 10 then  
      write_list (chars_of_Z (n + d));;  
      write (Int.repr newline);;  
      c <- read;;  
      read_sum (n + d) (Int.unsigned c - char0)  
    else ret tt else ret tt.
```

But really finite !

```
Definition main_itree := c <- read;; read_sum 0  
(Int.unsigned c - char0).
```

Definition main_spec :=

DECLARE _main WITH gv : globals

PRE [] main_pre_ext prog main_itree nil gv

POST [tint] main_post prog nil gv.

main_post == TT, so no ITREE here

I-Trees in VST: getchar/putchar

7. Add VST boilerplate and verify individual functions

Definition of Vprog, Gprog Espec...

Lemma `body_print_int`: `semax_body Vprog Gprog f_print_int print_int_spec`.

Lemma `body_print_intr`: `semax_body Vprog Gprog f_print_intr print_intr_spec`.

Lemma `body_main`: `semax_body Vprog Gprog f_main main_spec`.

I-Trees in VST: getchar/putchar

7. Add VST boilerplate and verify individual functions

Definition Vprog, Gprog Espec...

Not so fast...

Lemma body_print_int: semax_body Vprog Gprog f_print_int print_int_spec.
Lemma body_print_intr: semax_body Vprog Gprog f_print_intr print_intr_spec.
Lemma body_main: semax_body Vprog Gprog f_main main_spec.

Better definition of **ITREE**(t)

VST/veric/juicy_extspec.v

Definition **has_ext** {Z:Type} (ora:Z) :=

iospecs.v

Definition **ITREE** (t) := EX t' : _, !!(EquivUpToTau t t') && has_ext t'.

Remove tau actions (yields equiv relation)

Sanity checks:

Lemma **has_ext_ITREE** t: has_ext t |-- ITREE t.

Lemma **ITREE_impl** t t': EquivUpToTau t t' -> ITREE t |-- ITREE t'.

In particular:

- monad laws (associativity / congruence of bind) can be lifted to ITREE
- application-level equivalences can be exploited:

Lemma **write_list_app** l1 l2:
EquivUpToTau (write_list (l1 ++ l2)) (write_list l1;; write_list l2).

Better definition of **ITREE**(t)

VST/veric/juicy_extspec.v

Definition **has_ext** {Z:Type} (ora:Z) :=

iospecs.v

Definition **ITREE** (t) := EX t' : _, !!(EquivUpToTau t t') && has_ext t'.

Enjoy!

Sanity checks:

Lemma **has_ext_ITREE** t: has_ext t |-- ITREE t.

Lemma **ITREE_impl** t t': EquivUpToTau t t' -> ITREE t |-- ITREE t'.

In particular:

- monad laws (associativity / congruence of bind) can be lifted to ITREE
- application-level equivalences can be exploited:

Lemma **write_list_app** l1 l2:
EquivUpToTau (write_list (l1 ++ l2)) (write_list l1;; write_list l2).

I-Trees in VST: getchar/putchar

8. Putting it all together

Definition of Vprog, Gprog Espec...

Lemma `body_print_int`: `semax_body Vprog Gprog f_print_int print_int_spec`.

Lemma `body_print_intr`: `semax_body Vprog Gprog f_print_intr print_intr_spec`.

Lemma `body_main`: `semax_body Vprog Gprog f_main main_spec`.

Instance `Espec` : OracleKind :=
 add_funspecs IO_Espec (ext_link_prog prog) Gprog.

Theorem `prog_correct`: `semax_prog_ext prog main_itree Vprog Gprog`.

“`main_itree` is safe for
the whole `prog`”

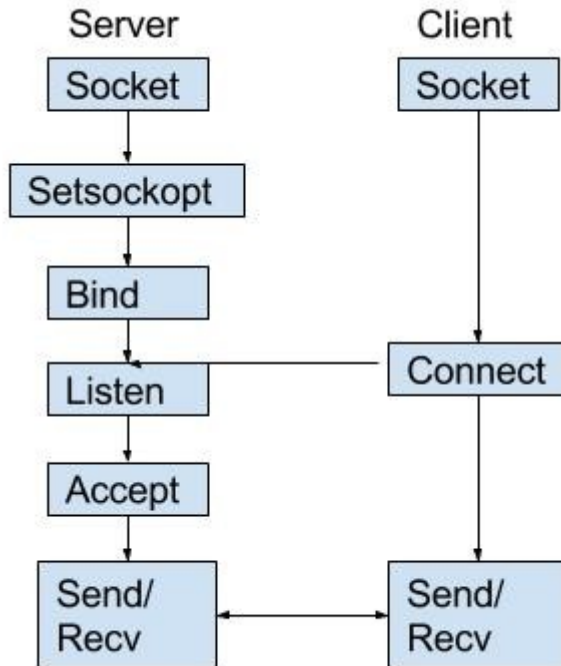
Connecting VST-Itrees to CertiKOS

Goal: show that CertiKOS indeed operates as specified, i.e. is a suitable "external word" for those ghost-state manipulating instructions that amount to OS calls.

Challenges:

- VST's ghost-supporting mathematical model not directly translatable into CertiKOS - concepts ("step-indexing", "predicates-in-the heap")
- But: erasing / compiling away step-indexing yields first-order specifications that amount to pre-post specifications over **CompCert** memories, these **are** understood by CertiKOS
- Current work: prove formal correspondence between step-erased specs and their ghost-enriched counterparts
- Approach restricted to OS calls without complex recursive structure (callbacks via function pointers), but sufficient for charaction-IO and network interactions: data sent/received is typically first-order

I-Trees in VST: sockets



<https://www.geeksforgeeks.org/socket-programming-cc/>

Function prototypes (server)

```
// returns sockfd IPv4 / IPv6... UDP / TCP Typically 0
int socket(domain, type, protocol);

// associates socket to address/port
int bind(int sockfd,
        const struct sockaddr *addr, socklen_t addrlen);

// put socket into "awaiting requests" mode
// Max length of queue of pending conn's
int listen(int sockfd, int backlog);

// establish conn with first request; return if of fresh conn
int accept(int sockfd,
          struct sockaddr *addr, socklen_t *addrlen);
```

Additional function prototype (client):

```
// try to establish new connection between sockfd and server at addr; retval indicates success
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

I-Trees in VST: sockets

1. Definition of I-Trees for socket interactions

Network-visible events:

Definition `connection_id` : Type := nat.

Definition `endpoint_id` : Type := nat.

Inductive `networkE` : Type -> Type :=

| `Listen` : endpoint_id -> networkE unit

| `Accept` : endpoint_id -> networkE connection_id

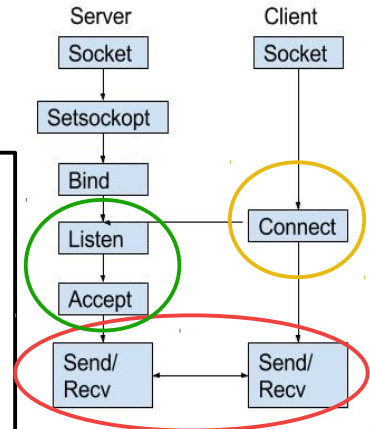
| `ConnectTo` : endpoint_id -> networkE connection_id

| `Shutdown` : connection_id -> networkE unit

| `RecvByte` : connection_id -> networkE (option string)

| `SendByte` : connection_id -> string -> networkE unit.

Lib/NetworkInterface.v



<https://www.geeksforgeeks.org/socket-programming-cc/>

I-Trees in VST: sockets

1. Definition of I-Trees for socket interactions

Derived Forms (excerpt):

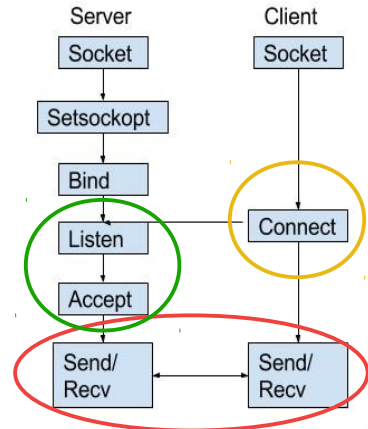
```
// Receive a string of length at most [len]. The return value [None] signals that  
// a connection was closed, when modelling the [recv()] POSIX syscall.
```

Definition `recv` $\{\text{networkE} \rightarrow E\} \{\text{nondetE} \rightarrow E\}$
(`c` : `connection_id`) (`len` : `nat`) : `M E (option bytes)` :=

```
//Send all bytes in a bytestring.
```

Fixpoint `send` $\{\text{networkE} \rightarrow E\}$ (`c` : `connection_id`)
(`bs` : `bytes`) : `M E unit` := `for_bytes bs (send_byte c)`.

Lib/NetworkInterface.v:



<https://www.geeksforgeeks.org/socket-programming-cc/>

Extend network events by failure + nondeterminism:

Definition `SocketE` : `Type` -> `Type` := (`nondetE` +' `failureE` +' `networkE`).

I-Trees in VST: sockets

1. Definition of I-Trees for socket interactions

Protocol stages of socket: these actions are **not** network-visible (OS-interaction only)

Lib/Socketinterface.v:

Inductive `socket_status` :=

| `ClosedSocket`
| `OpenedSocket`
| `BoundSocket` (addr : endpoint_id)
| `ListeningSocket` (addr : endpoint_id)
| `ConnectedSocket` (conn : connection_id).

again, client-side...

{ z:Z | 0 <= z <= maxFD }

Record `SocketMap` : **Type** := { lookup_socket: sockfd -> socket_status }.

We capture that status of all sockets in a second mpred:

Parameter `SOCKAPI`: SocketMap -> mpred.

Will soon be replaced by another ghost variable...

I-Trees in VST: sockets

2a. ITREE predicate

Again, don't use `has_ext` directly but

Definition `SocketMonad` := M SocketE.

Definition `ITREE` {T} (t : SocketMonad T) :=
EX t' : SocketMonad T, `!!(trace_incl t t') && has_ext t'`.
Spec/Vst/SocketSpecs.v

where

Definition `trace_incl` {R : Type} (m m' : M_R) : Prop :=
forall t r, `is_trace_of t r m -> is_trace_of t r m'`.

Lib/Socketinterface.v

Inductive `is_trace_of` {R : Type}: trace -> option R -> M_R -> Prop := ...

Lib/Socketinterface.v

generalizes `equivUpToTau` from previous example, yielding, e.g.

Lemma `trace_bind_assoc`: forall ...,

Spec/Vst/SocketSpecs.v

`ITREE (b <- (a <- m ;; f a) ;; g b) = ITREE (a <- m ;; b <- f a ;; g b)`.

I-Trees in VST: sockets

2b. VST axioms for send / receive

```
Definition recv_spec (T : Type) := Spec/Vst/SocketSpecs.v
  DECLARE _recv

  WITH t : SocketMonad T, k : option string -> SocketMonad T, st : SocketMap,
    client_conn : connection_id, fd: sockfd, buf_ptr: val, alloc_len: Z, sh: share

  PRE [ 1%positive OF tint, 2%positive OF (tptr tvoid), 3%positive OF tuint, 4%positive OF tint ]

  PROP ( consistent_world st; lookup_socket st fd = ConnectedSocket client_conn; writable_share sh;
    trace_incl (r <- recv client_conn (Z.to_nat alloc_len) ;; k r) t )

  LOCAL (temp 1%positive (Vint (Int.repr (descriptor fd))); temp 2%positive buf_ptr;
    temp 3%positive (Vint (Int.repr alloc_len)); temp 4%positive (Vint (Int.repr 0)))

  SEP ( SOCKAPI st; ITREE t;
    data_at_sh (Tarray tuchar alloc_len noattr) buf_ptr )

  POST [ tint ] EX result : _, st' : SocketMap, r : _, contents : _

  PROP ( 0 <= r <= alloc_len ∨ r = NO;
    r > 0 → st'=st ∧ ∃! msg, result = inr (Some msg) ∧ Zlength (val_of_string msg) = r ∧
      sublist 0 r contents = (val_of_string msg) ∧
      sublist r alloc_len contents = list_repeat (Z.to_nat (alloc_len - r)) Vundef;
    (*case EOF*) r = 0 -> result = inr None ∧ contents = list_repeat (Z.to_nat alloc_len) Vundef ∧
      st' = update_socket_state st fd OpenedSocket;
    r < 0 -> st'=st ∧ result = inl tt ∧ contents = list_repeat (Z.to_nat alloc_len) Vundef;
    Zlength contents = alloc_len; consistent_world st' )

  LOCAL ( temp ret_temp (Vint (Int.repr r)) )

  SEP ( data_at sh (tarray tuchar alloc_len) contents buf_ptr; SOCKAPI st' ;
    ITREE (match result with
      | inl tt => t
      | inr msg_opt => k msg_opt
    end)).
```

I-Trees in VST: sockets

2b. VST axioms for send / receive

```
Definition recv_spec (T : Type) := Spec/Vst/SocketSpecs.v
  DECLARE _recv

  WITH t : SocketMonad T, k : option string -> SocketMonad T, st : SocketMap,
    client_conn : connection_id, fd: sockfd, buf_ptr: val, alloc_len: Z, sh: share

  PRE [ 1%positive OF tint, 2%positive OF (tptr tvoid), 3%positive OF tuint, 4%positive OF tint ]

  PROP ( consistent_world st; lookup_socket st fd = ConnectedSocket client_conn; writable_share sh;
    trace_incl (r <- recv client_conn (Z.to_nat alloc_len) ;; k r) t )
```

Yes, the spec for send is similar (and similarly lengthy)...

**Specs for bind and socket don't involve
ITREE terms, just SOCKAPI terms.**

Again, these axioms need to be justified by CertiKOS.

```
r < 0 -> st = st' /\ result = inl tt /\ contents = list_repeat (Z.to_nat alloc_len) vunder,
Zlength contents = alloc_len; consistent_world st')

LOCAL ( temp ret_temp (Vint (Int.repr r)) )

SEP ( data_at sh (tarray tuchar alloc_len) contents buf_ptr; SOCKAPI st' ;
  ITREE (match result with
    | inl tt => t
    | inr msg_opt => k msg_opt
  end)).
```

I-Trees in VST: sockets

3. Write some C program: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c`

I-Trees in VST: sockets

3. Write some C program: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c`

... 631 LOC



I-Trees in VST: sockets

3. Write some C program: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c`

... 631 LOC



4. Write an informal spec: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c` – **lines 2- 10!**

I-Trees in VST: sockets

3. Write some C program: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c`

... 631 LOC



4. Write an informal spec: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c` – lines 2- 10!

This program implements a swap server that accepts connections from multiple clients.

On each connection, the server receives bytes from the corresponding client until the buffer is full, at which point the received bytes are deemed a "message".

The server replies to this client with the last obtained message, which could have been obtained from another client (on another connection). This last obtained message is initially set to all zeroes.

All these actions (accepting a new connection, receiving a message, and sending a message) are interleaved.

I-Trees in VST: sockets

3. Write some C program: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c`

... 631 LOC



4. Write an informal spec: `dw/DeepWeb/Spec/main.cw/DeepWeb/Spec/main.c` – lines 2- 10!

This program implements a swap server that accepts connections from multiple clients.

On each connection, the server receives bytes from the corresponding client until the buffer is full, at which point the received bytes are deemed a "message".

The server replies to this client with the last obtained message, which could have been obtained from another client (on another connection). This last obtained message is initially set to all zeroes.

All these actions (accepting a new connection, receiving a message, and sending a message) are interleaved.



I-Trees in VST: sockets

5. Translate the informal spec into Coq / Itrees

“ClikeSpec” : counterparts of data structures

```
Inductive connection_state : Type :=
  RECVING | SENDING | DONE |
  DELETED.
```

```
Record connection : Type :=
{ conn_id : connection_id;
  conn_request : string;
  conn_response : string;
  conn_response_bytes_sent : Z;
  conn_state : connection_state
}.
```

```
typedef enum state {
  RECVING,
  SENDING,
  DONE,
  DELETED,
} state;
```

```
typedef struct connection {
  socket_fd fd;
  uint32_t request_len;
  uint8_t request_buffer[BUFFER_SIZE];
  uint32_t response_len;
  uint8_t response_buffer[BUFFER_SIZE];
  uint32_t num_bytes_sent;
  enum state st;
  struct connection* next;
} connection;
```

I-Trees in VST: sockets

5. Translate the informal spec into Coq / Itrees

“ClikeSpec” : counterparts of data structures

Inductive `connection_state` : Type :=
 RECVING | SENDING | DONE |
 DELETED.

Record `connection` : Type :=
 { conn_id : connection_id;
 conn_request : string;
 conn_response : string;
 conn_response_bytes_sent : Z;
 conn_state : connection_state
 }.

```
typedef enum state {
  RECVING,
  SENDING,
  DONE,
  DELETED,
} state;
```

```
typedef struct connection {
  socket_fd fd;
  uint32_t request_len;
  uint8_t request_buffer[BUFFER_SIZE];
  uint32_t response_len;
  uint8_t response_buffer[BUFFER_SIZE];
  uint32_t num_bytes_sent;
  enum state st;
  struct connection* next;
} connection;
```

and their operations

Definition `upd_conn_request` (conn : connection) (request : string) : connection := ...

Definition `upd_conn_response` (conn : connection) (response : string) : connection := ...

etc.

I-Trees in VST: sockets

5. Translate the informal spec into Coq / Itrees

counterparts of C functions

(* Wait for a message from connection [conn]. [recv] can return a partial message, in which case we store the bytes we received to try receiving more in a later iteration. *)

Definition `conn_read` (buffer_size : Z) (conn: connection) (last_full_msg : string)

: **M socketE** (connection * string) := . . .



```
static int
conn_read(connection* conn, store* last_msg_store) {
    . . .
}
```

I-Trees in VST: sockets

5. Translate the informal spec into Coq / Itrees

counterparts of C functions

(* Wait for a message from connection [conn]. [recv] can return a partial message, in which case we store the bytes we received to try receiving more in a later iteration. *)

Definition `conn_read` (buffer_size : Z) (conn: connection) (last_full_msg : string)

: M socketE (connection * string) := . . .

28 lines ; -)



20 lines

```
static int
conn_read(connection* conn, store* last_msg_store) {
    . . .
}
```

. . . similarly for the other data structures and functions of the C code . . .

I-Trees in VST: sockets

6. Embed Itrees in VST specs

```

Definition conn_read_spec (T : Type) (buffer_size : Z) :=
  DECLARE _conn_read
  WITH k : (connection * string) -> SocketM T, st : SocketMap, conn : connection,
    fd : sockfd, last_msg : string, conn_ptr : val, msg_store_ptr : val
  PRE [ _conn OF (tptr (Tstruct _connection noattr)), _last_msg_store OF (tptr (Tstruct _store noattr)) ]
  PROP ( consistent_world st; conn_state conn = RECVING ;
    consistent_state buffer_size st (conn, fd) )
  LOCAL ( temp _conn conn_ptr ; temp _last_msg_store msg_store_ptr )
  SEP ( SOCKAPI st ; ITREE (r <- conn_read buffer_size conn last_msg ;; k r) ;
    list_cell LS Tsh (rep_connection conn fd) conn_ptr ;
    field_at Tsh (Tstruct _store noattr) [] (rep_store last_msg) msg_store_ptr )
  POST [ tint ] EX conn' : _, EX last_msg' : _, EX st' : _, EX r : Z,
  PROP ( r = YES ; consistent_world st' ;
    recv_step buffer_size (conn, fd, st, last_msg) (conn', fd, st', last_msg') )
  LOCAL ( temp ret_temp (Vint (Int.repr r)) )
  SEP ( SOCKAPI st' ; ITREE (k (conn', last_msg')) ;
    list_cell LS Tsh (rep_connection conn' fd) conn_ptr ;
    field_at Tsh (Tstruct _store noattr) [] (rep_store last_msg') msg_store_ptr ).
  
```

etc.

I-Trees in VST: sockets

7. Add VST boilerplate and verify individual functions

■ ■ ■ ■

Lemma `body_conn_read`:

`semax_body Vprog Gprog f_conn_read (conn_read_spec unit BUFFER_SIZE).`

Proof. `start_function. ... Qed.`

430 lines

etc.

I-Trees in VST: sockets

8. Putting it all together

dw/DeepWeb/Proofs/Vst/verif_main.v

Lemma `body_main`: `semax_body Vprog Gprog f_main (main_spec server)`.

Proof. ... `Qed`.

dw/DeepWeb/Proofs/Vst/verif_prog.v

Lemma `all_funcs_correct`: `semax_func Vprog Gprog (prog_func prog) Gprog`.

Proof. ... `Qed`.

Swap server: top level claim / proof

“**tree** is safe for the whole **prog**”

Spec/TopLevelSpec.v

Definition swap_server_correct := exists (**tree** : M socketE unit),
semax_prog_ext **prog tree** Vprog Gprog \wedge
refines_mod_network swap_observer (simplify_network **tree**)

Currently omits listen, shutdown

“**Refinement scrambling**”: any scrambling that results from a trace of the **low-level implementation** can also arise from a trace of the **high-level spec**.

Swap server: top level claim / proof

“**tree** is safe for the whole **prog**”

Spec/TopLevelSpec.v

Definition `swap_server_correct` := exists (**tree** : M socketE unit),
semax_prog_ext **prog tree** Vprog Gprog \wedge
`refines_mod_network` `swap_observer` (simplify_network **tree**)

Currently omits listen, shutdown

“**Refinement scrambling**”: any scrambling that results from a trace of the **low-level implementation** can also arise from a trace of the **high-level spec**.

Proofs/TopLevelProof.v

Theorem `swap_server_is_correct` : swap_server_correct.
Proof. exists **main_itree**; split. ... **Admitted**.

(*ToDo: scaffolding macros, parts of **refinement** *)

Swap server: top level claim / proof

“**tree** is safe for the whole **prog**”

Spec/TopLevelSpec.v

```
Definition swap_server_correct := exists (tree : M socketE unit),  
  semax_prog_ext prog tree Vprog Gprog /\  
  refines_mod_network swap_observer (simplify_network tree)
```

Currently omits listen, shutdown

Notation:

“ prog \models swap_observer “

“Refinement so
results from a trace
can also arise from a trace of the high-level spec.”


Proofs/TopLevelProof.v

```
Theorem swap_server_is_correct : swap_server_correct.  
Proof. exists main_itree; split. ... Admitted.
```

(*all_funcs_correct verified but some scaffolding / macros ToDo *)


Scaling out (in all dimensions)

Horizontally:




Scaling out (in all dimensions)

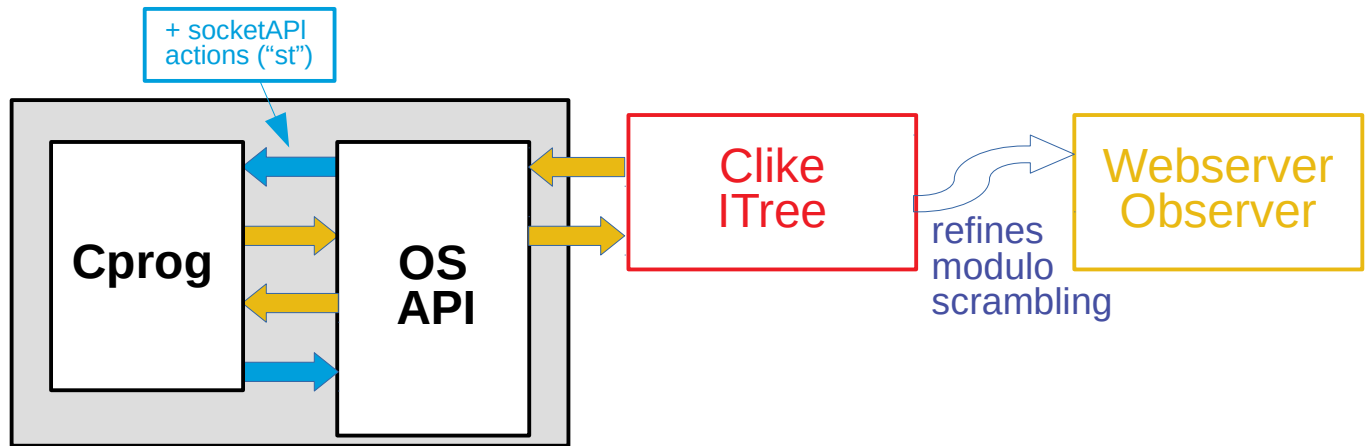
Horizontally:



Vertically:




or better




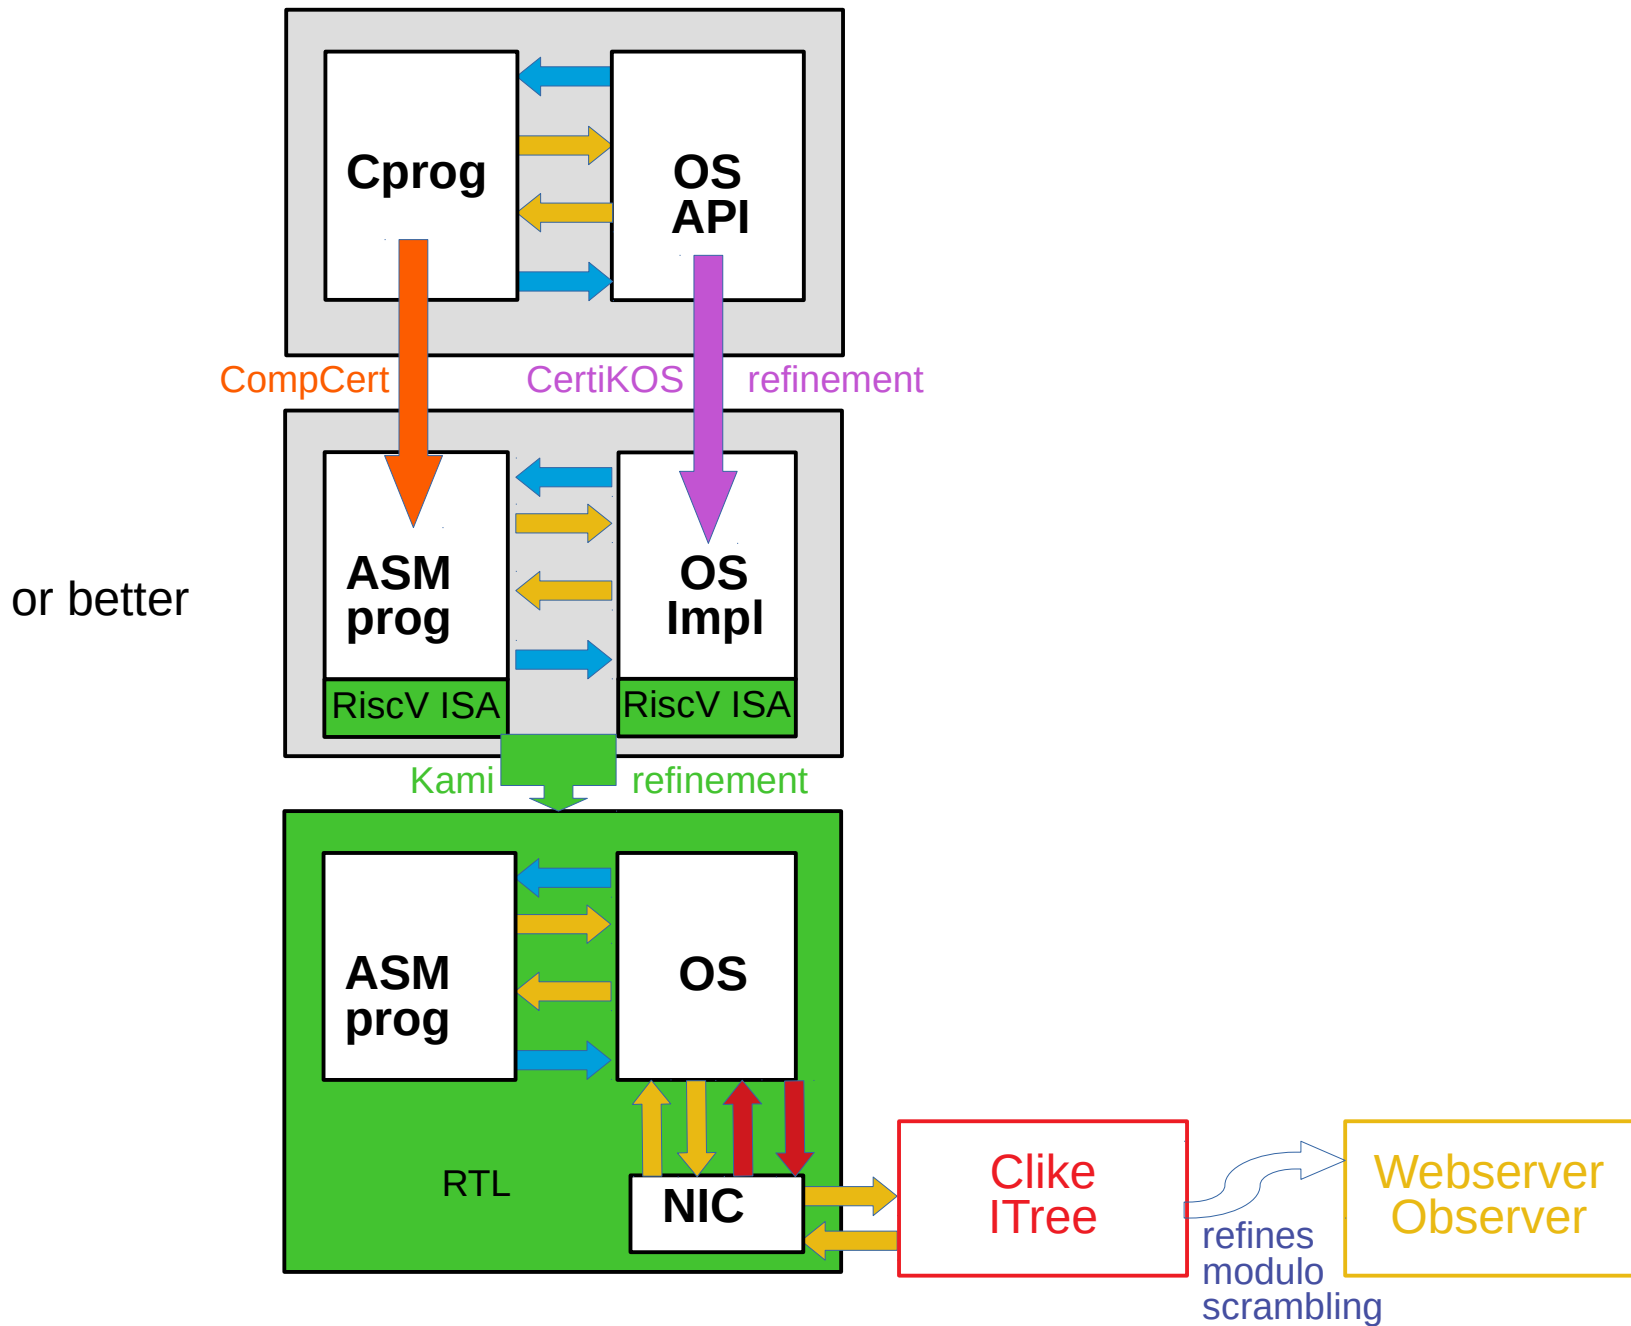
Scaling out (in all dimensions)

Horizontally:




Vertically:





Challenge I: composition of various refinement notions



Challenge I: composition of various refinement notions

Challenge II:

Theorem Correct:

Cprog \models server.

Proof. ... Qed.



Integration testbed for DeepSpec Technologies

